



AWS Managed Services DynamoDB

Adriano Sastre Vieira
adrianosastre@inatel.br
adrianosastre

AWS MANAGED SERVICES DYNAMODB

AGENDA

AWS - The 5 Pillars

Serverless / Microservices

AWS Managed Services

Minimum Stack:

API Gateway + Lambda + DynamoDB

DynamoDB - Theory

What is it / What is not

How to model (NoSql vs RDBMS)

Single Table Design

When/how to use

When/how NOT to use

DynamoDB - Hands on

Cloud Formation + CDK

The project on Visual Code

Deploying on AWS


Testing with Postman

AWS - THE 5 PILLARS

<https://aws.amazon.com/getting-started/fundamentals-core-concepts/>

New "mental models":

- **Operational excellence**: thinking about operations as automation (CloudFormation, CDK ...)
- **Security**: zero trust, the principle of least privilege (IAM, data encryption ...)
- **Reliability**: using fault isolation zones to limit blast radius (regions, quotas, throttling ...)
- **Performance Efficiency**: think of your services as cattle, not pets. (horizontal vs vertical scaling, auto scaling ...)
- **Cost Optimization**: OpEx (pay-as-you-go) instead of CapEx (one-time purchase)



SOME IMPORTANT CONCEPTS

Serverless:

- Serverless is the practice of using **managed services** with **event driven** compute functions to avoid or minimize infrastructure management, configuration, operations, and idle capacity.
- Wide range of things an application architecture may need. e.g. **Compute, Storage, Data, Monitoring, Queue, Notification** ...
- It's a paradigm shift, potentially even more impactful than the move to *the "Cloud"* before it

FaaS: Function as a Service

- Serverless approach to **compute** (i.e.: AWS Lambda)
- But compute (i.e. running code) is not the sole aspect of the application architecture

<https://medium.com/serverless-transformation/in-defence-of-serverless-the-term-764514653ea7>

<https://pauldjohnston.medium.com/serverless-best-practices-b3c97d551535>

SOME IMPORTANT CONCEPTS

Microservices:

- The microservice **architectural style** is an approach to developing a single application as a **suite of small services**, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API.
- These services are built around **business capabilities** and **independently deployable** by fully automated deployment machinery.
- Many people are using **serverless applications** to build a **microservice architecture**.

<https://martinfowler.com/articles/microservices.html>

Advantages:

- **Independence** and **decoupling** of parts of the system
- You can choose the most appropriate **technologies / database** per service
- Small teams with the whole knowledge in each service is better for development and maintenance

Main characteristic:

- One database per microservice

Challenges:


- Distributed transactions between services
- Design for failure
- Interface changes needs to be well coordinated between services

SOME AWS MANAGED SERVICES

<https://aws.amazon.com/serverless/>

<https://serverlessland.com/>

<https://aws.amazon.com/free>




PROS:

- Cost reduction
 - initial cost tend to **zero** (free tier, pay as you use)
- Security
 - e.g. no need to install security patches on servers
- High **performance**
- NoOps
 - e.g. no servers / storage to manage or to scale, no need for a infrastructure team
- More **scalable**
- Greener
- **Productivity**: developers **focus** on delivering **business value**

SERVERLESS EVENT-DRIVEN DESIGN SAMPLE ARCHITECTURE

(STARTING WITH THE BASIC)



<https://aws.amazon.com/getting-started/deep-dive-serverless/>

API GATEWAY

Fully **managed service** that makes it easy for developers to **create, publish, maintain, monitor,** and **secure** APIs at any **scale**.

Handles all the tasks involved in accepting and processing up to hundreds of thousands of concurrent API calls, including **traffic management, CORS support, authorization and access control, throttling, monitoring,** and **API version** management.

REST APIs

WebSocket APIs

Free-tier 12-months applied, up to 1 million requests per month, after that: cost by request and volume of data transferred.



<https://aws.amazon.com/api-gateway>

<https://aws.amazon.com/free/?all-free-tier.q=api-gateway>

LAMBDA

09

AWS Lambda is an **event-driven**, serverless **computing** platform provided by Amazon as a part of Amazon Web Services.

Runs code in response to events and **automatically manages the computing resources** required by that code.

- Triggered by events (e.g. HTTP Calls via API Gateway, S3 new objects on a bucket, new SQS in a queue, new item in a DynamoDB table ...)

Free-tier forever applied, up to 1 million requests / month; after that: cost by allocated memory and execution time consumed.



<https://aws.amazon.com/lambda>

<https://aws.amazon.com/free/?all-free-tier.q=lambda>

DYNAMODB

Fast and flexible NoSQL database service for any scale. Key-value and document database that delivers single-digit millisecond performance at any scale.

Fully managed, multi-region, multi-active, durable database with built-in security, backup and restore, and in-memory caching for internet-scale applications.

Free-tier forever applied, up to 25GB storage, 25 read/write provisioned units (about 200M requests/month). After that: cost by storage and requests.

<https://www.dynamodbguide.com/the-dynamo-paper>

- Created by Amazon (2004-2007), public released by AWS in 2012

<https://aws.amazon.com/dynamodb/>

<https://aws.amazon.com/free/?all-free-tier.q=dynamoDB>



DYNAMODB AS PART OF A SERVERLESS ARCHITECTURE



11

Just one part of moving to a more serverless architecture often involves adopting [DynamoDB](#) in place of the relational databases of the past.

NoSQL isn't new, but the idea of using it for any and all core business requirements is — this is because the whole system now works at a [scale and event-driven nature](#) before unknown.

[DynamoDB](#) is complex, doing it well even more so.

[https://medium.com/
serverless-
transformation/in-
defence-of-serverless-
the-term-
764514653ea7](https://medium.com/serverless-transformation/in-defence-of-serverless-the-term-764514653ea7)

DYNAMODB NOSQL DESIGN VS RDBMS



12

NoSQL design requires a **different mindset** than RDBMS design.

RDBMS = you can go ahead and create a **normalized** data model without thinking about **access patterns**.

By contrast, **you shouldn't start designing your schema for DynamoDB until you know the questions it will need to answer**. Understanding the business problems and the application use cases up front is essential.

Access Patterns, e.g. :

- Get a user's profile data
- List the user's orders
- Get an order and its items
- List the user's orders by status

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-general-nosql-design.html#bp-general-nosql-design-concepts>

DYNAMODB DESIGN CONSIDERATIONS

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-general-nosql-design.html#bp-general-nosql-design-concepts>



13

It is important to understand three fundamental properties of your application's access patterns before you begin:

*"DynamoDB is a **Key-value and document database that delivers single-digit millisecond performance at any scale.**"*

- Data **size**: Knowing how much data will be stored and requested at one time will help determine the most effective way to **partition** the data.
 - Data **velocity**: DynamoDB scales by increasing the number of physical partitions that are available to process queries, and by efficiently distributing data across those partitions. Knowing in advance what the peak query loads will be might help determine how to partition data to best use I/O capacity.
 - Data **shape**: Instead of reshaping data when a query is processed (as an RDBMS system does), a NoSQL database organizes data so that **its shape in the database corresponds with what will be queried**. This is a key factor in increasing **speed** and **scalability**.
- **size / velocity**:
 - Is it really needed? (such a scale and performance)
 - Does it make sense to use a NoSQL database? Is it the best option?
 - Is the technical team familiar with it?
 - **shape**
 - Do you know all the questions that needs to be answered in the system?
 - **NO (for any of the above)**
 - Have you considered using a RDS datase, e.g. **Amazon Aurora Serverless**?
- YES we are going to use DynamoDB!
- Then data **shape** is the most important aspect!

DYNAMODB DESIGN

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-general-nosql-design.html#bp-general-nosql-design-concepts>



14

After you identify specific query requirements, you can **organize data** according to general **principles** that govern **performance**:

- **Keep related data together.** Keeping related data in close proximity has a major **impact** on **cost** and **performance**. Instead of distributing related data items across multiple tables, you should keep related items in your NoSQL system as close together as possible. **As a general rule, you should maintain as few tables as possible in a DynamoDB application.**
- **Use sort order.** Related items can be grouped together and queried efficiently if their **key design** causes them to **sort together**. This is an important NoSQL design strategy.
- **Distribute queries.** It is also important that a high volume of queries not be focused on one part of the database, where they can exceed I/O capacity. Instead, you should design data keys to **distribute traffic evenly across partitions as much as possible**, avoiding "hot spots."
- **Use global secondary indexes.** By creating specific global secondary indexes, you can **enable different queries** than your main table can support, and that are still **fast** and **relatively inexpensive**.

DYNAMODB DESIGN

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-general-nosql-design.html#bp-general-nosql-design-concepts>



15

After you identify specific query requirements, you can organize data according to general principles that govern performance:

- Keep related data together. Keeping related data in close proximity has a major impact on cost and performance. Instead of distributing related data items across multiple tables, you should keep related items in your NoSQL system as close together as possible. As a general rule, you should maintain as few tables as possible in a DynamoDB application.
- Use sort order. Related items can be grouped together and queried efficiently if their key design causes them to sort together. This is an important NoSQL design strategy.
- Distribute queries. It is also important that a high volume of queries not be focused on one part of the database, where they can exceed I/O capacity. Instead, you should design data keys to distribute traffic evenly across partitions as much as possible, avoiding "hot spots."
- Use global secondary indexes. By creating specific global secondary indexes, you can enable different queries than your main table can support, and that are still fast and relatively inexpensive.

These general principles translate into some common design patterns that you can use to model data efficiently in DynamoDB.

DYNAMODB

PK, SK, ATTRIBUTES



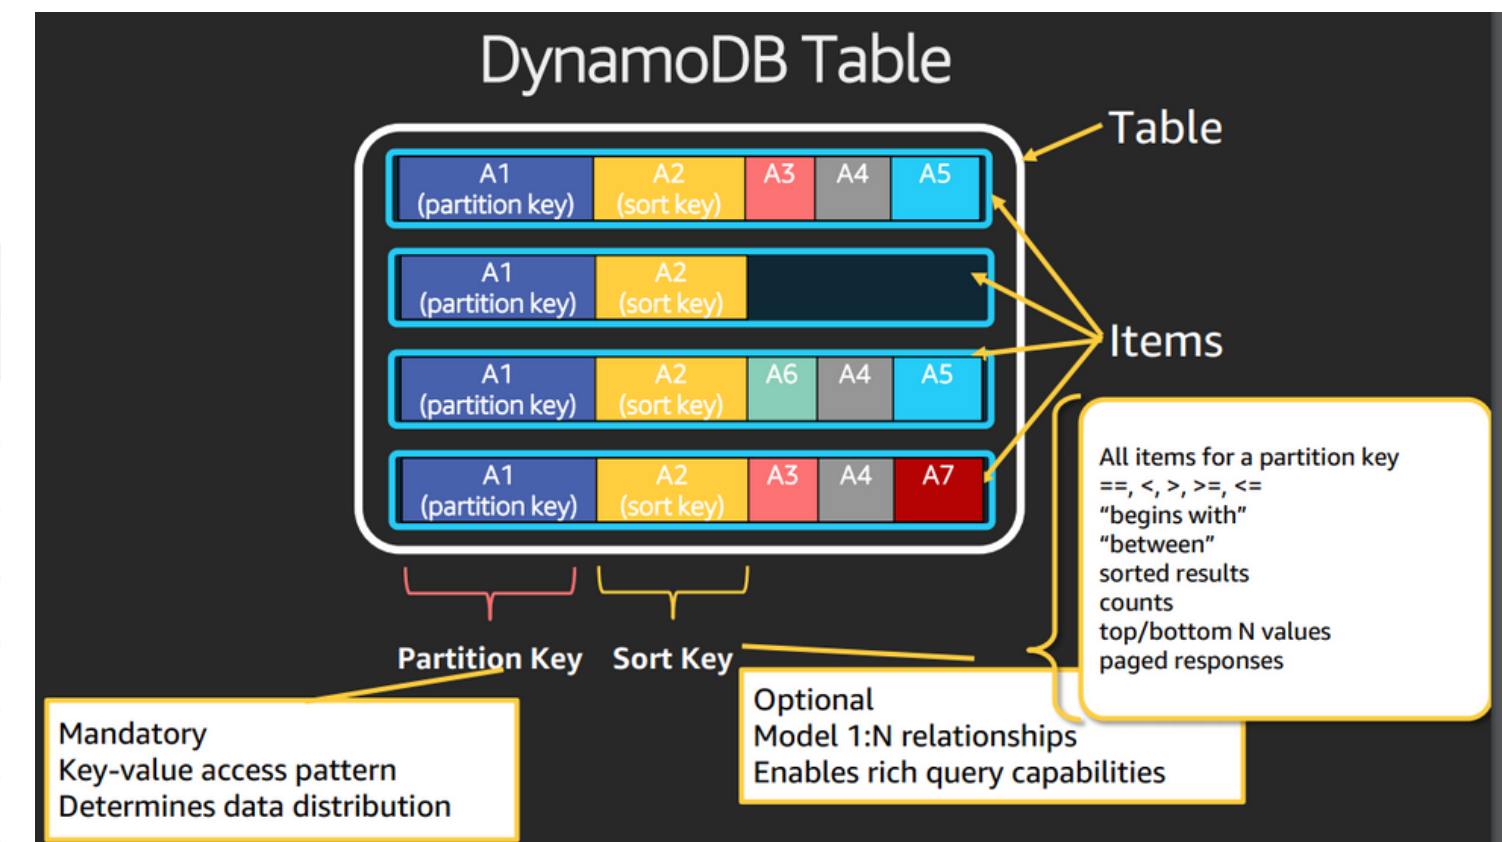
16

- Primary Key:
 - PK = **Partition** Key (hash key) (mandatory)
 - SK = **Sort** Key (range key) (optional)
- **Attributes**: Binary, Number or String
 - Can be grouped in a JSON-like structure

Each table item may have different attributes!

<https://www.youtube.com/watch?v=kSnpuKr3Ajw>

Primary Key		Attributes		
Actor (PARTITION)	Movie (SORT)			
Tom Hanks	Cast Away	Role	Year	Genre
		Chuck Noland	2000	Drama
Tim Allen	Toy Story	Role	Year	Genre
		Woody	1995	Children's
Natalie Portman	Toy Story	Role	Year	Genre
		Buzz Lightyear	1995	Children's
	Black Swan	Role	Year	Genre
		Nina Sayers	2010	Drama



DYNAMODB CAPACITY



17

Per table configuration.

OnDemand:

- automatic and "infinite" scaling

Provisioned:

- possible to define independent read and write unit capacities
- possible to configure auto-scaling rules, e.g. min/max scaling
- when correctly configured, less cost than onDemand

Possible to change between OnDemand and Provisioned 1x per day!

<https://aws.amazon.com/dynamodb/pricing/>

DYNAMODB

HOW TO QUERY DATA



18

Get = one specific item, by PK or PK/SK


Query = several items, by PK/SK or indexed attributes

- High performance

Scan = several items, by any table attribute

- Poor performance
- High cost

DYNAMODB INDEXES



19

GSI = Global Secondary Index = applied to all table items

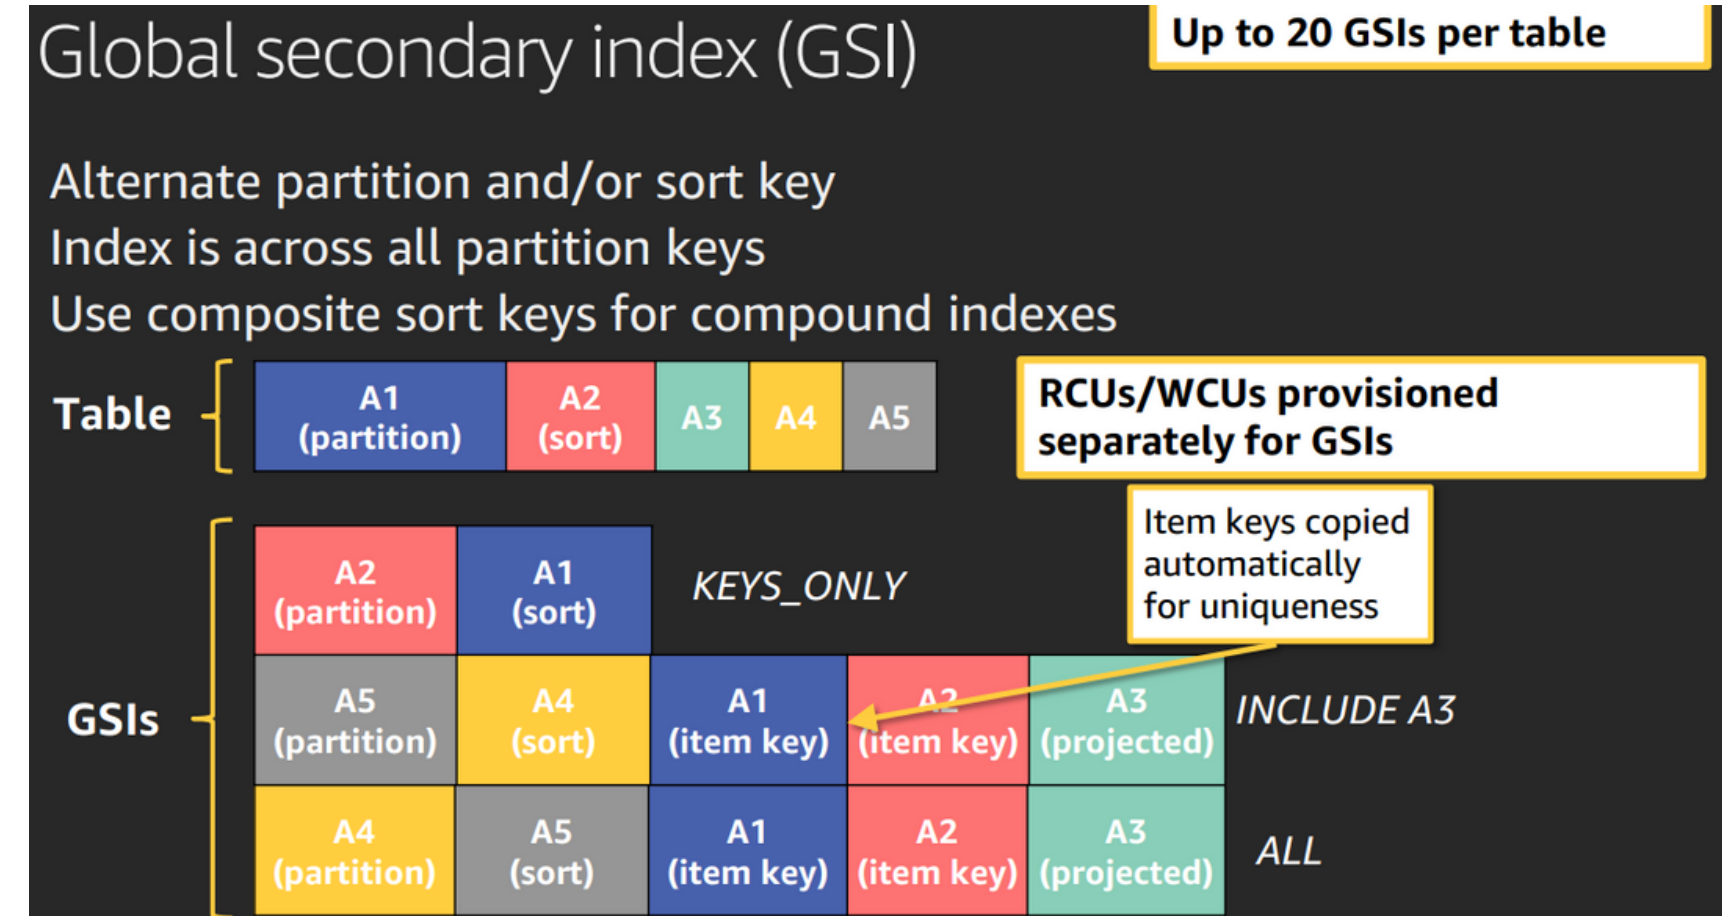
LSI = Local Secondary Index = Index inside the table partition (PK)

https://d1.awsstatic.com/events/reinvent/2019/R/EPEAT1_Advanced_design_patterns_for_Amazon_DynamoDB_DAT334-R1.pdf

Pros: performance: with an index, it is possible to query (instead of scan) on attributes other than the PK/SK

Cons: behind the scenes, each GSI duplicates the table storage (and costs)

Limit = 20 GSI per table!



DYNAMODB

TTL



20

TTL (time to live attribute)


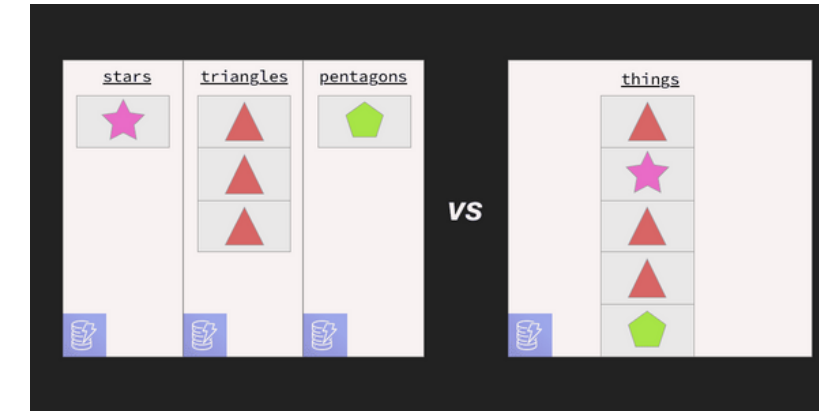
Optional attribute.

Automatically deletes the table item when the TTL is reached.

Each item may have a different time to live!

DYNAMODB

SINGLE TABLE DESIGN



21

AWS Recommendation: maintain as few tables as possible!


Relational databases = recommend to **normalize** the data

Pro: data access very flexible

Con: it reduces the scalability (joins)

DynamoDB was built for **enormous**, **high-velocity** use cases, such as the Amazon.com shopping cart.

Rather than working to make joins scale better, DynamoDB sidesteps the problem by **removing the ability to use joins at all**.



DYNAMODB

SINGLE TABLE DESIGN

DynamoDB was build with **web scale** in mind.
It can grow almost infinitely without degrading performance.

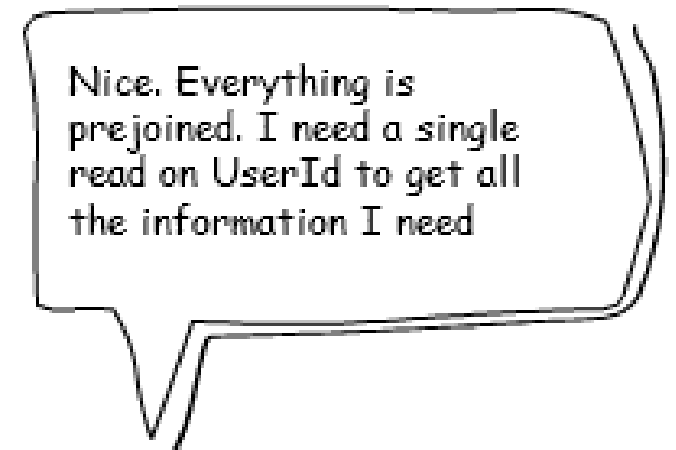
To achieve this DynamoDB removed joins completely.

You have to model the data in such a way that you can read the data in a **single request** by **denormalizing** the data.


*The main reason for using a single table in DynamoDB is to retrieve **multiple, heterogenous** item types using a **single request**.*



22



Single Table		
UserId	UseInfo	
	Order1	
	Order2	



Order History			
<input type="checkbox"/>	=	<input type="checkbox"/>	=
<input type="checkbox"/>	=	<input type="checkbox"/>	=
<input type="checkbox"/>	=	<input type="checkbox"/>	=

DYNAMODB

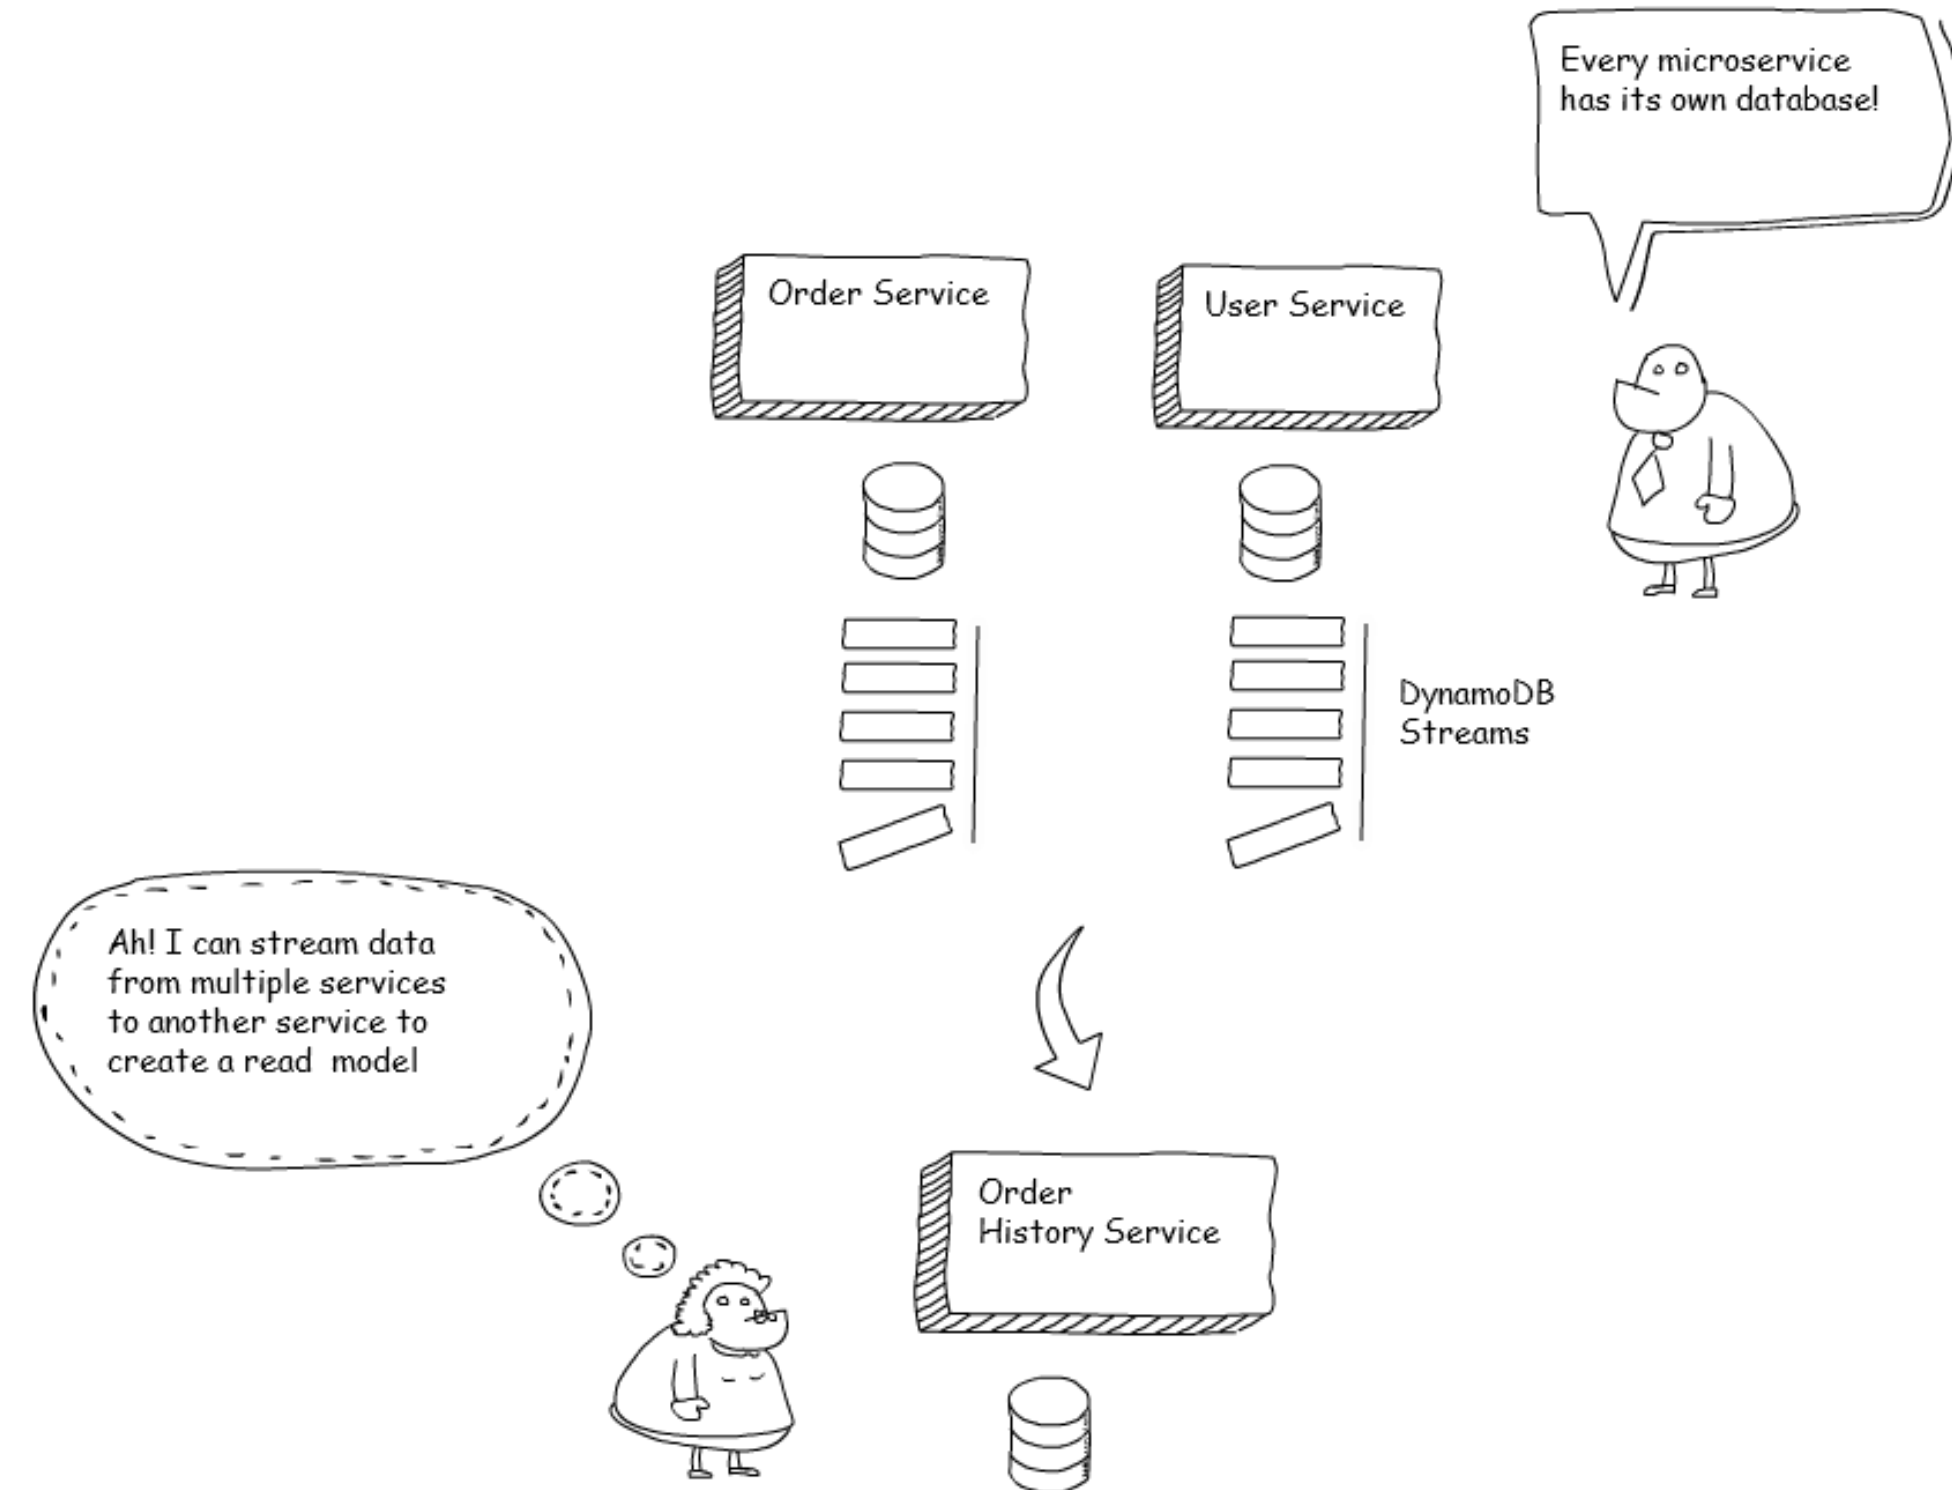
SINGLE TABLE DESIGN VS MICROSERVICES



23

One table per service vs One single table for all microservices?

The best practice for microservices is that every microservice owns its own data.



DYNAMODB

SINGLE TABLE DESIGN VS MICROSERVICES



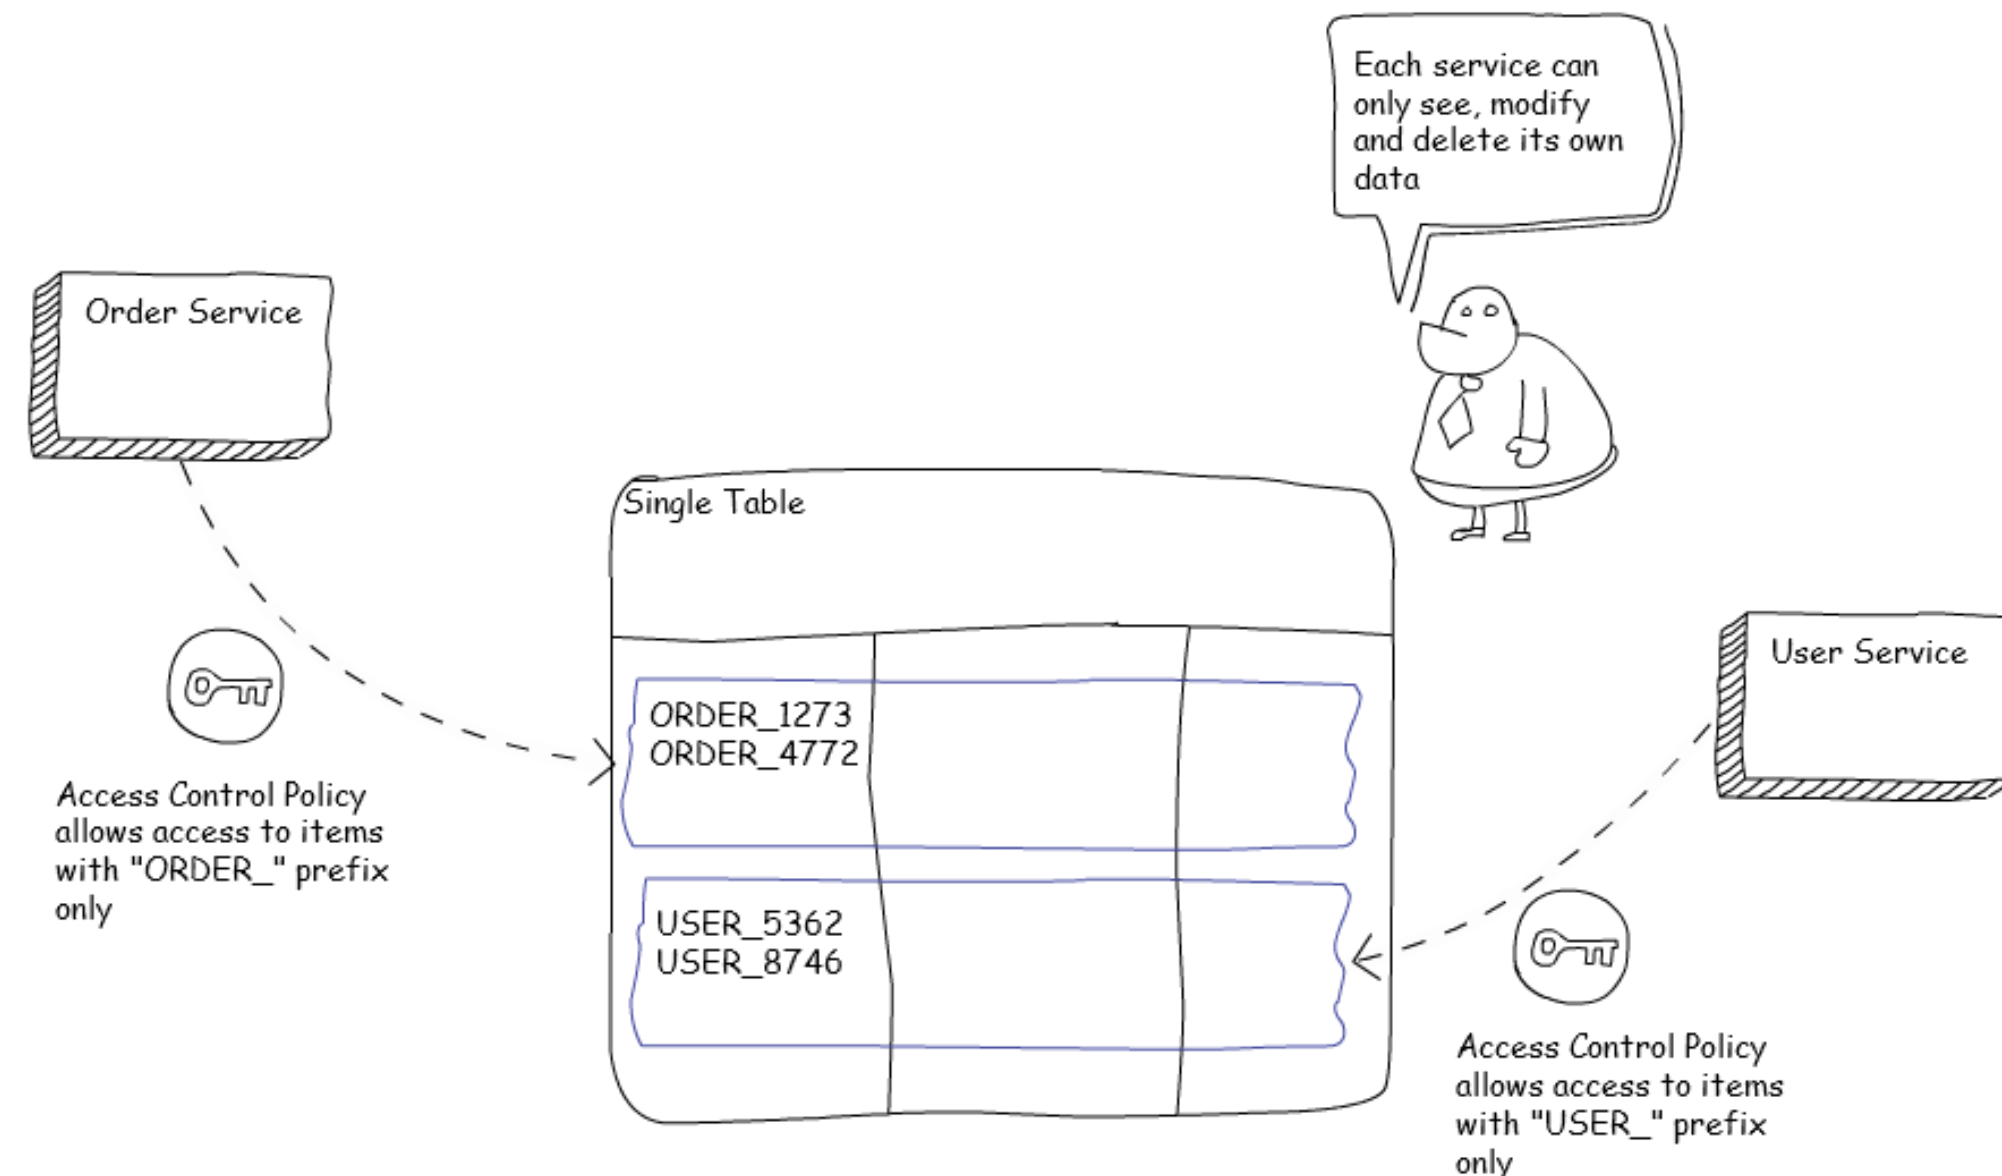
24

Alternative: use a single table with namespaces

HOW =

1. **prefix** every primary key (partition or sort key) with the name of the service
2. add a **fine grained access control** policy to your service that allows access to items with the service's prefix only.

<https://cbannes.medium.com/microservices-with-dynamodb-should-you-use-a-single-table-or-use-one-table-per-microservice-25f54cf610d9>



DYNAMODB

SINGLE TABLE DESIGN

PROS X CONS



25

- Reducing the number of requests (Performance!)
- Less operational overhead (alarms, metrics...)
- Save money (capacity, provisioning for read /write ...)
- When well designed, less indexes in total (e.g. a "username" GSI can be reused)
- Steep learning curve
- **Inflexibility of adding new access patterns**
- Difficulty of exporting your tables for analytics
- *"NoSQL is not flexible. It's efficient but not flexible." — Rick Houlihan*

<https://www.alexdebrie.com/posts/dynamodb-single-table/>
<https://serverlessfirst.com/dynamodb-modelling-single-vs-multi-table/>
<https://www.jeremydaly.com/how-to-switch-from-rdbms-to-dynamodb-in-20-easy-steps/>

"A well-optimized single-table DynamoDB layout looks more like machine code than a simple spreadsheet"

DYNAMODB

WHEN AND HOW TO USE



26

When you have any of the following:

- **Huge** amount of data
- **Performance** required
- Well defined data **access patterns**

<https://www.alexdebrie.com/posts/dynamodb-single-table/>

<https://www.alexdebrie.com/posts/dynamodb-patterns-serverless/#the-true-microservice>

*When modeling with DynamoDB, you should be following best practices. This includes **denormalization**, **single-table design**, and other proper **NoSQL modeling** principles.*

→ In this pattern, the **True Microservice**, we will see how to get all of the benefits of DynamoDB with none of the downsides.

DYNAMODB

WHEN AND HOW NOT TO USE !



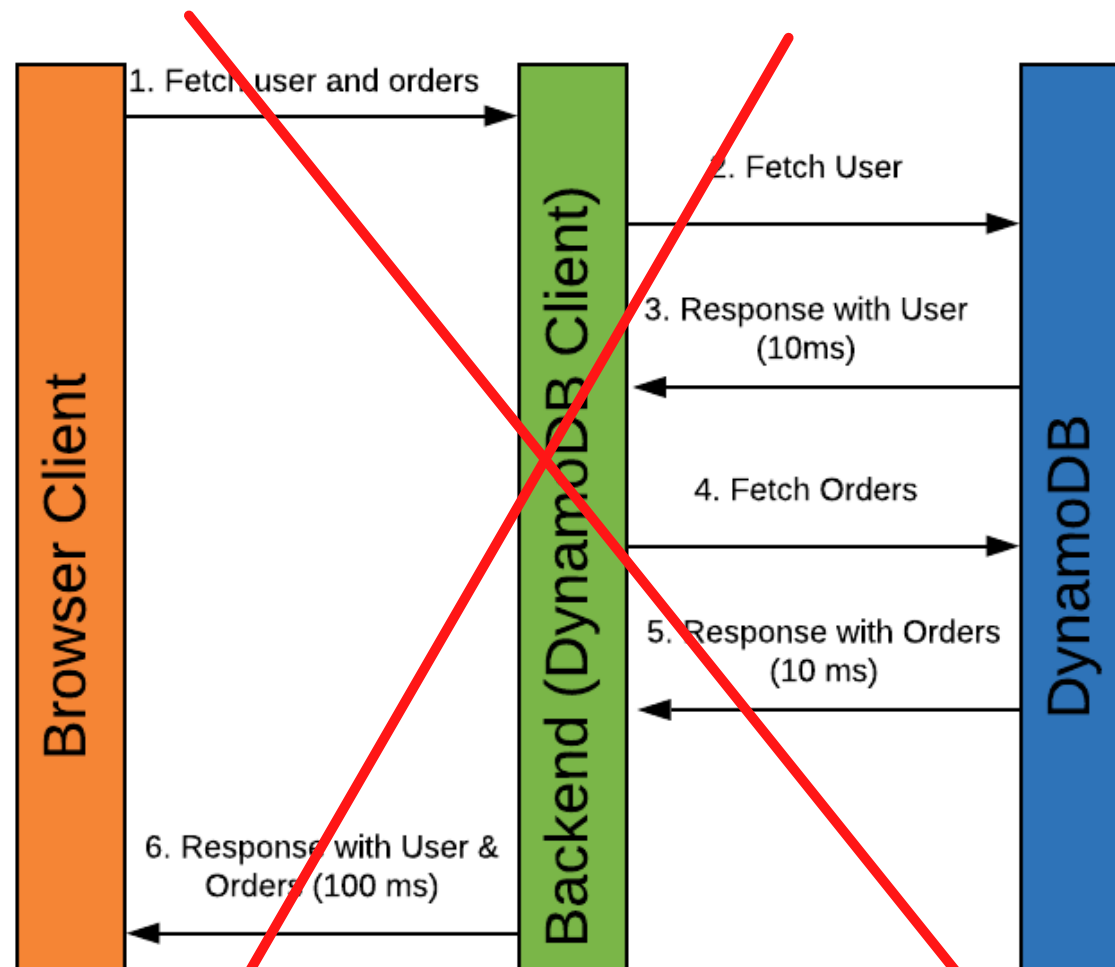
27

- “Whenever the benefits don’t outweigh the costs”

- In new applications where developer agility is more important than application performance
- In applications using GraphQL

Faux-SQL design: the practice of choosing parts of both relational and NoSQL concepts, e.g. use multiple DynamoDB tables and normalize their data.

- This pattern is **not** recommended by the DynamoDB team
- This pattern recommended a lot by the AWS AppSync team (!)



- You are making **multiple queries** for your calls. You should aim to satisfy your needs with a single query.
- Your **application** code just became a **database management system**. Rather than doing a JOIN in SQL, you are doing it in **code**. This means you are **responsible** for maintaining **referential integrity** and for **optimizing** your queries.

You know what’s really good at maintaining referential integrity and optimizing JOINS? **A relational database**.

<https://www.alexdebrie.com/posts/dynamodb-patterns-serverless/#faux-sql>

HANDS ON


The image features a white background with two red geometric shapes in the corners. In the top-left corner, a red triangle is partially visible, with a black line extending from its vertex towards the center. In the bottom-right corner, a red triangle is also partially visible, with a black line extending from its vertex towards the center.

Thank you!

adrianosastre@inatel.br

HANDS ON

08



This project on github:

https://github.com/adrianosastre/XPForward_DynamoDB_CDK